

Summer 2023

---

New York Institute of Technology  
College of Engineering and Computing Sciences

# Regression Modelling for House Price Predictions

DTSC 701/CSCI 636 Big Data Analytics  
Professor Liangwen Wu

Stevenson Chittumuri (1226890), Albert Chen (1299225)

Due Date: August 25th, 2023

**Abstract**—This research project sits at the intersection of cloud computing, data architecture, and predictive modeling using California’s housing data to predict housing prices. Our main objective is to run our regression models to ascertain trends in the data set that may provide insight into California’s housing price distribution and second, to guide cloud service utilization and data storage optimization. We have provided a step-by-step guide to run the Spark application on AWS. Once up and running, three regression models were deployed in a production AWS environment: Linear Regression, Decision Tree Regression, and Random Forest Regression. Out of the set of regression models implemented, Random Forest Regression emerged as the most promising, with approximately 75.63% of the target variance being explained by the predictors and the least average prediction error. The study concludes by running the California Housing Data Set on a KNN Regression Model not available in our current production environment to offer more insight into the trends in the data set that require the integration of domain knowledge.

## CONTENTS

<b>I</b>	<b>Introduction</b>	<b>2</b>
<b>II</b>	<b>Data Storage Decision: Data Lake vs. Data Warehouse</b>	<b>2</b>
<b>III</b>	<b>Documentation for Running Spark Application on AWS</b>	<b>3</b>
<b>IV</b>	<b>Data Visualization</b>	<b>5</b>
<b>V</b>	<b>Regression Modeling Evaluation</b>	<b>6</b>
<b>VI</b>	<b>Future Improvements</b>	<b>7</b>
	<b>Appendix A: Regression Modelling Code</b>	<b>9</b>
	<b>References</b>	<b>11</b>

## I. INTRODUCTION

This project report aims to elucidate the complexities of cloud computing, data architecture decisions, and model development and evaluation, through data analysis and regression modeling of California’s housing Data Set [2]. This report provides not just insights into the state’s housing price distribution, but also offers a framework and guide for effectively leveraging cloud services and optimizing data storage choices. Our choice of the California Housing Data Set is intentional given its full-featured data set with low dimensionality allowing ease of usage in regression modeling deployed in a cloud production environment with data storage capabilities.

## II. DATA STORAGE DECISION: DATA LAKE VS. DATA WAREHOUSE

When choosing a storage system, our decision came down to using a data lake or a data warehouse. AWS, our chosen cloud service, offers robust solutions for both. For data lakes, AWS provides Amazon S3 (Simple Storage Service), which is designed to store and retrieve any amount of data. S3 is reliable, scalable, and offers easy data access. Given its capabilities, it was an appealing choice. On the data warehouse side, AWS offers Amazon Redshift, a fully managed, petabyte-scale data warehouse service. Redshift is optimized for online analytic processing (OLAP) and business intelligence tools.

While Redshift offers fast querying capabilities, it's primarily structured for this type of workload.

We specifically chose the data lake approach using Amazon S3 for our project. Due to its scalability, Amazon S3 easily scales to accommodate growing data without any upfront costs [3]. In terms of flexibility, S3 can store both structured and unstructured data. Given the diverse data types in our data set, this was a significant advantage. With regards to cost-effectiveness, Amazon S3 uses a pay-as-you-go model aligned well with our project's budget and goals. While Amazon Redshift is a powerful data warehouse option, we prioritized flexibility and scalability over purely fast querying capabilities for this particular project. Thus, the data lake approach using Amazon S3 was more in line with our needs.

### III. DOCUMENTATION FOR RUNNING SPARK APPLICATION ON AWS

**Initial Setup:** Create an AWS account and have CSV and Python scripts downloaded.

#### Step 1: Setting Up Your EC2 Instance

- Log in to your AWS Management Console.
- Navigate to EMR on EC2 and 'Create Cluster'.
- Choose the 'Spark' application, and 'm5.xlarge' instance types for Primary, Core, and Tasks.
- Provision '3' for the instance sizes for Primary and Task.
- Find "Security configuration and EC2 key pair", and be sure to create or select an SSH Key Pair to remotely access the cluster from your machine. Be

sure to have it downloaded onto your machine.

- Finally, create or select roles under the 'Identity and Access Management (IAM) roles' section.

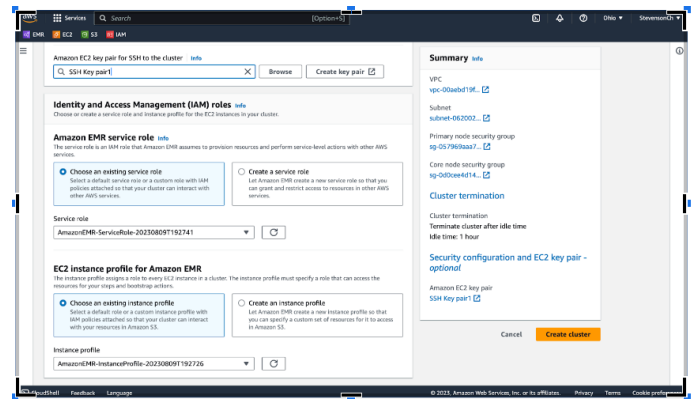


Fig. 1. Shows that there are no missing/misconfigured values on the right side, allowing me to create the cluster.

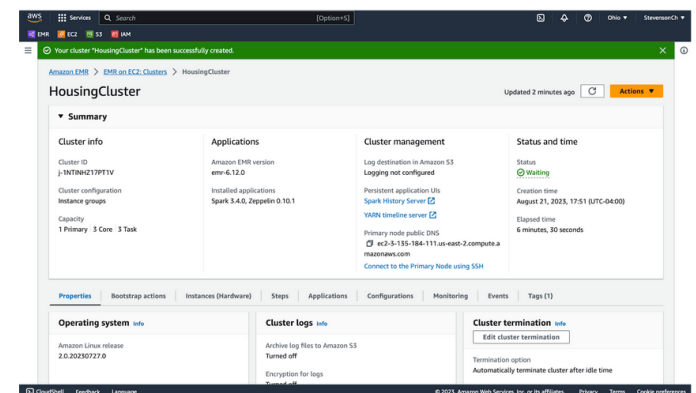


Fig. 2. Successfully created an online cluster with a public DNS node available for SSH

#### Step 2: Enable SSH privileges to EC2 Cluster

- When you open the cluster, scroll down to the "EC2 security groups (firewall)" section and open the hyperlink for your primary node's security group.
- Inside the "Inbound rules" section be sure to add SSH (port 22) privileges for your IP address.
- If you don't have SSH privileges, click on "Edit inbound rules", scroll down to "Add rule"

and choose SSH and the desired IPV4 address (preferably “0.0.0.0/0”) for global access as this is just a test. Save rules and you should see “SSH” as a rule in your primary group.

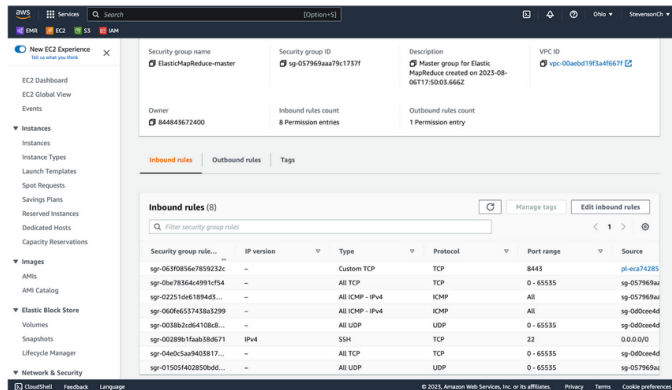


Fig. 3. SSH protocol is part of the inbound rules now

### Step 3: Upload dataset to S3 Bucket (data lake)

- In AWS Console, open S3 and create a bucket, and uncheck “Block all public access” (if you don’t already have one dedicated to this project).
- Then inside this bucket, click on ‘Upload’ and upload the CSV file containing the dataset.
- After uploading, select the file and click on ‘Copy S3 URI’ and paste that link into the Python script in the read CSV method.

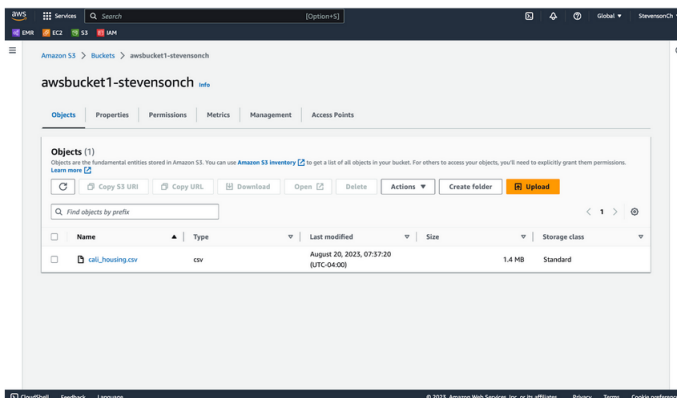


Fig. 4. File inside the data lake S3 bucket

### Step 4: Upload Your Script to the Cluster

- Open the terminal (or command prompt) on your machine
- Use the ‘scp’ command to securely copy your Python script to the EC2 instance:

```
scp -i /path/to/your/ssh_key.pem /path/to/spark_cali_housing.py hadoop@your-ec2-ip-address:/home/hadoop/
```

```
prince -- -zsh -- 117x24
(base) prince@Princes-MBP ~ % scp -i /Users/prince/Desktop/SSH/Key/ pair1.pem /Users/prince/Documents/DTSC_Project/s
park_cali_housing.py hadoop@ec2-3-135-184-111.us-east-2.compute.amazonaws.com:/home/hadoop/
park_cali_housing.py
(base) prince@Princes-MBP ~ %
```

Result in terminal if ran successfully

Fig. 5. Result in terminal if ran successfully

### Step 5: SSH into Your EC2 Instance

- Inside your terminal session, we will access your EC2 instance using the ‘ssh’ command:

```
ssh -i /path/to/your/ssh_key.pem hadoop@your-ec2-ip-address
```

Fig. 6. SSH command

- After you are SSH’ed into the cluster, type ‘ls’ to see if the Python script has been uploaded.

### Step 6: Run Your Spark Application on the Cluster

- Execute your Spark application using the ‘spark-submit’ command:

During execution, you should see Spark start processes and allocate resources to spark-related



on the correlation rates in the column for median house value, it is evident that median income has the highest positive correlation with median house value with a correlation rate of 77%. The next highest correlation rate for a feature in the data set is ocean proximity with a correlation rate of 27%.




Fig. 12. Correlation Matrix

The high positive correlation between median income and median house value is highlighted in the scatter plot below. The strong positive relationship between the median income and house value will have a strong effect on our training models.

## V. REGRESSION MODELING EVALUATION

In running a range of regression models on the California Housing Data Set, we obtained the following results for Linear Regression, Decision Tree Regression, and Random Forest Regression. The Performance of each training model is summarized in R-squared and Root Mean Squared Error [1].

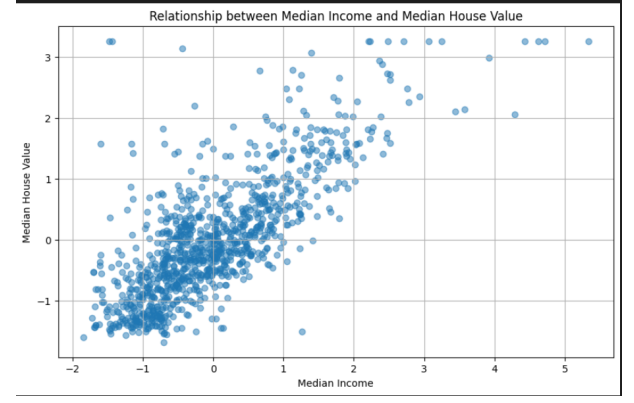


Fig. 13. Scatter plot: relation between median income and median house value

R-squared measures the proportion of the variance in the output variable based on the input variables. A perfect fit between the feature set and the label set would yield a value of 1. With values ranging from 0 to 1, R-squared communicates the amount of noise between each feature used to predict the label. An analysis of the results of our training models resulted in Random Forest Regression having the highest R-square value of 75.64%, while Linear Regression and Decision Tree Regression had similar R-squared statistics of 62.68% and 62.81%, respectively.

Regression Modelling Performance Metrics			
	Linear Regression	Decision Tree Regression	Random Forest Regression
R-squared	0.6268	0.6079	0.7563
Root Mean Square Error	0.6531	0.6694	0.5278

Fig. 14. Performance Metrics

In addition to R-squared we also analyzed performance through Root Mean Squared Error. The classic metric used in gradient descent for regression modeling, Root Mean Squared Error is a loss function that measures the difference between predicted and actual values passed to the regression

model. Given that a lower Root Mean Squared Error value is optimal indicating a lesser cost between the prediction and actual value, Random Forest Regression not only offers the best fit between the feature set and label set but is also the most accurate with a Root Mean Squared Error of 0.5278. This is significantly more accurate than the Root Mean Squared Error of Linear Regression and Decision Tree Regression which have error rates of 0.6531 and 0.6519 units, respectively. Overall, the Random Forest Regression is the optimal regression model out of the set of regression models used on the California Housing Data Set. The combination of a relatively high R-squared value and a low Root Mean Squared Error suggests greater accuracy in its predictive powers.

## VI. FUTURE IMPROVEMENTS

For our project, we implemented three supervised learning algorithms: Linear Regression, Decision Tree Regression, and Random Forest Regression. To run further experiments we could also expand our training models to include an unsupervised learning algorithm, such as KNN Regression to derive further insights from our data sets. As Decision Tree Regression and Linear Regression perform similarly a trend in the data cannot be ascertained between the two training models. In fact, in running the KNN Regression outside of our AWS cluster we were able to obtain the following results for running a

KNN Regression Model on the California Housing Data Set used in our initial three supervised learning trials.

Regression Modelling Performance Metrics				
	Linear Regression	Decision Tree Regression	KNN Regression	Random Forest Regression
R-squared	0.6268	0.6079	<b>0.6845</b>	0.7563
Root Mean Square Error	0.6531	0.6694	<b>0.6005</b>	0.5278

Fig. 15. KNN Regression Metrics

From the figure above it is evident that KNN Regression is second to Random Forest Regression in both R-squared and Root Mean Square Error metrics, with a 68.45% R-squared and 60.05% Root Mean Square Error Rate. The predictive qualities based on the feature set are second highest out of all the regression models used according to the R-square metric, while the error rate captured in the Root Mean Square Error metric is the second lowest. In both instances, KNN Regression comes in second to Random Forest Regression in terms of performance indicating a relatively high accuracy in comparison to the rest of the regression models implemented in the project.




Fig. 16. KNN Regression Scatter Plot

From the strong performance results of KNN Regression and the associated KNN Regression Scatter Plot (Fig. 16), it is evident that data localization

occurs within the data set, as certain home owner's with similar income levels tend to buy homes with similar housing prices in the state of California. Given the unsupervised nature of KNN Regression it is also important to confirm our findings with a domain expert. Due to the limitations of the MLlib Spark API, our project did not include the KNN Regression Model when connecting to our AWS cluster. A future iteration of the project could integrate Spark's latest API in order to incorporate a KNN Regression Model with the set of regression models already deployed to a production environment.



## APPENDIX A

### REGRESSION MODELLING CODE

```

1 from pyspark.sql import SparkSession
2 from pyspark.ml.feature import StandardScaler,
  VectorAssembler, Imputer, StringIndexer
3 from pyspark.ml.regression import LinearRegression,
  RandomForestRegressor, DecisionTreeRegressor
4 from pyspark.ml.evaluation import
  RegressionEvaluator
5 from pyspark.ml import Pipeline
6 import pyspark.sql.functions as F
7 from itertools import chain
8
9 # Create Spark session
10 spark = SparkSession.builder.appName("housing").
  getOrCreate()
11
12 # Load data from AWS S3
13 housing_df = spark.read.csv("s3://awsbucket-achen33/
  housing.csv", header=True, inferSchema=True)
14
15 # Assuming the following mapping for '
  ocean_proximity' (you can change this
  accordingly)
16 myMapCat = {
17     'NEAR BAY': 1,
18     'INLAND': 2,
19     '<1H OCEAN': 3,
20     'NEAR OCEAN': 4,
21     'ISLAND': 5
22 }
23
24 map_expr = F.create_map([F.lit(x) for x in chain(*
  myMapCat.items())])
25 housing_df = housing_df.withColumn("ocean_proximity"
  , map_expr[housing_df["ocean_proximity"]])
26
27 # Handle Missing Values
28 imputer = Imputer(strategy="median", inputCols=["
  total_bedrooms"], outputCols=["total_bedrooms"])
29 housing_df = imputer.fit(housing_df).transform(
  housing_df)
30
31 # Feature Scaling
32 features = [col_name for col_name in housing_df.

```

```

  columns if col_name != "median_house_value" and
  col_name != "ocean_proximity"]
33 assembler = VectorAssembler(inputCols=features,
  outputCol="features")
34 scaler = StandardScaler(inputCol="features",
  outputCol="scaled_features")
35
36 # Train-Test Split
37 train_data, test_data = housing_df.randomSplit([0.7,
  0.3], seed=0)
38
39 # Linear Regression
40 lin_reg = LinearRegression(featuresCol="
  scaled_features", labelCol="median_house_value")
41 pipeline_lin = Pipeline(stages=[assembler, scaler,
  lin_reg])
42 model_lin = pipeline_lin.fit(train_data)
43 predictions_lin = model_lin.transform(test_data)
44
45 # Decision Tree Regression
46 dt_reg = DecisionTreeRegressor(featuresCol="
  scaled_features", labelCol="median_house_value")
47 pipeline_dt = Pipeline(stages=[assembler, scaler,
  dt_reg])
48 model_dt = pipeline_dt.fit(train_data)
49 predictions_dt = model_dt.transform(test_data)
50
51 # RandomForest
52 rf_reg = RandomForestRegressor(featuresCol="
  scaled_features", labelCol="median_house_value",
  numTrees=100)
53 pipeline_rf = Pipeline(stages=[assembler, scaler,
  rf_reg])
54 model_rf = pipeline_rf.fit(train_data)
55 predictions_rf = model_rf.transform(test_data)
56
57 evaluator = RegressionEvaluator(labelCol="
  median_house_value", predictionCol="prediction")
58
59 # Evaluate Linear Regression
60 r2_lin = evaluator.evaluate(predictions_lin, {
  evaluator.metricName: "r2"})
61 rmse_lin = evaluator.evaluate(predictions_lin, {
  evaluator.metricName: "rmse"})
62 print("Linear Regression - R-squared:", r2_lin)
63 print("Linear Regression - Root Mean Squared Error:"

```

```
        , rmse_lin)

64
65 # Evaluate Decision Tree
66 r2_dt = evaluator.evaluate(predictions_dt, {
    evaluator.metricName: "r2"})
67 rmse_dt = evaluator.evaluate(predictions_dt, {
    evaluator.metricName: "rmse"})
68 print("Decision Tree - R-squared:", r2_dt)
69 print("Decision Tree - Root Mean Squared Error:",
    rmse_dt)
70
71 # Evaluate Random Forest
72 r2_rf = evaluator.evaluate(predictions_rf, {
    evaluator.metricName: "r2"})
73 rmse_rf = evaluator.evaluate(predictions_rf, {
    evaluator.metricName: "rmse"})
74 print("Random Forest - R-squared:", r2_rf)
75 print("Random Forest - Root Mean Squared Error:",
    rmse_rf)
76
77 # Close Spark session
78 spark.stop()
```

## REFERENCES

- [1] Funmilola Alaba Ajala Ololade Oyewo Yetunde Faith Akande Gbenle Oluwadara Abigail Bola Adetunji, Oluwatobi Noah Akande. House price prediction using random forest machine learning technique. *Procedia Computer Science*, pages 806–813, 2022.
- [2] Cam Nugent. California housing prices.
- [3] Durai Raj Vincent Dharmendra Singh Rajput Shaik Abdul Khalandar Basha, Syed Muzamil Basha. Chapter 11 - challenges in storing and processing big data using hadoop and spark. *Academic Press*, pages 179–187, 2019.